Machine Learning examples in High Energy Physics

Oleg Filatov[1], Kirill Bukin[2]

[1]Deutsches Elektronen-Synchrotron (DESY)
[2]Moscow Institute of Physics and Technology

Overview

Overview

Disclaimer: it is impossible to cover all the applications at once. Those which are presented here have been selected solely for the sake of supporting the main narrative of the presentation. I encourage you to follow the links on the slides, explore what's not mentioned and see how far your curiousity may get you!

Note: since the beginning of my research path I've always been a member of CMS collaboration, hence a corresponding bias in presented applications might take place.

Making theories

Making theories

- Note: I am extremelly ”experiment”-biased and don’t really keep track of theory developments
- Here for the sake of completeness I highlight just a few applications which I stumbled upon – but there’s definitely more

Making theories

- Note: I am extremelly "experiment"-biased and don't really keep track of theory developments
- Here for the sake of completeness I highlight just a few applications which I stumbled upon – but there's definitely more
- Lattice QCD : accelerated sampling of field configuration with **gauge-equivariant** *normalising flows*

Approximate the QCD path integral by **Monte Carlo**

$$\langle \mathcal{O} \rangle = \frac{1}{Z} \int \mathcal{D}A \mathcal{D}\overline{\psi} \mathcal{D}\psi \mathcal{O}[A, \overline{\psi}\psi] e^{-S[A,\overline{\psi}\psi]} \rightarrow \langle \mathcal{O} \rangle \simeq \frac{1}{N_{\mathrm{conf}}} \sum_i^{N_{\mathrm{conf}}} \mathcal{O}([U^i])$$

with field configurations U^i distributed according to $e^{-S[U]}$

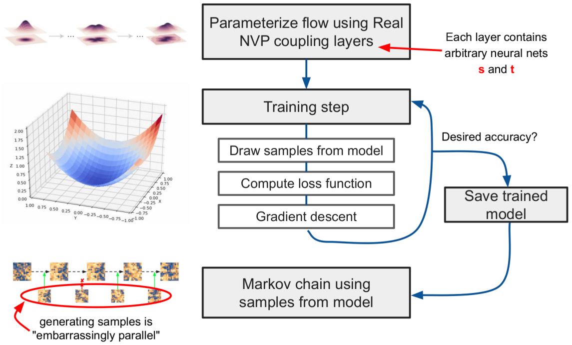

Summary chart: Tej Kanwar

Making theories

- Note: I am extremelly "experiment"-biased and don't really keep track of theory developments
- Here for the sake of completeness I highlight just a few applications which I stumbled upon – but there's definitely more
- Lattice QCD : accelerated sampling of field configuration with **gauge-equivariant** *normalising flows*
- String Theory: computation, optimisation, topology

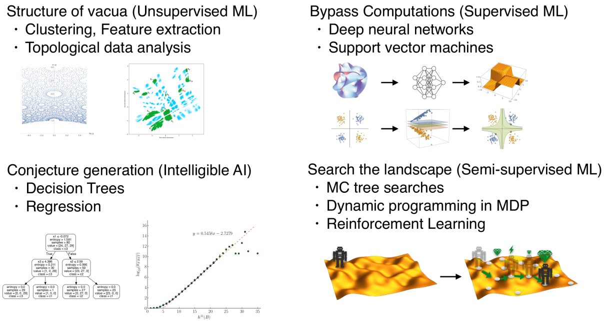

Building HEP setups

- It all starts from building a detector - which is tough, **a lot of constraints** need to be taken into account (money vs. physics)
- Can we somehow find the "optimal" geometry?

Building HEP setups

Detector geometry

- It all starts from building a detector - which is tough, **a lot of constraints** need to be taken into account (money vs. physics)
- Can we somehow find the "optimal" geometry?
- Yes, as it's an optimisation problem
- In fact, **"black box"** optimisation → bayesian optimisation

➔ Muon shielding in SHiP experiment

➔ LHCb ECal case

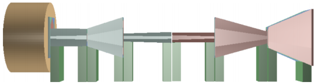

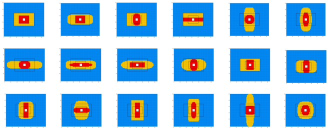

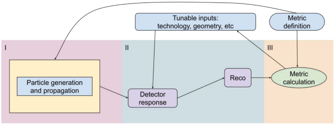

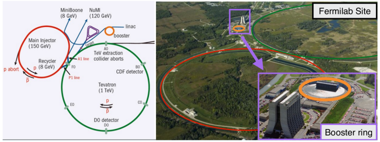

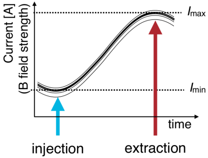

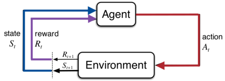

Booster: 400 MeV—> 8 GeV

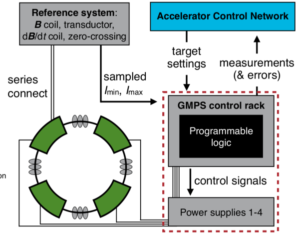

- Combination of RF cavities and bending magnets
- Bending magnet current ramps in 15hz cycles to maintain the orbit of the accelerating beam

finally, treat it with love and care

Collecting data

Ngadiuba's talk @FastML2020

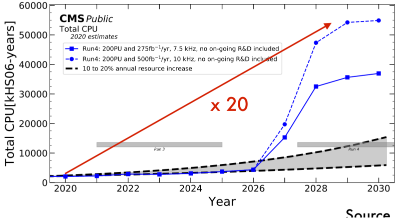

- Once collisions begin, data from the detector needs to be collected
- Huge in size *but* can't store everything
- Respond and process new data quickly *but* computational resources limited

→ Need to find the "best", select it and do this fast → requires **low latency, large throughput**

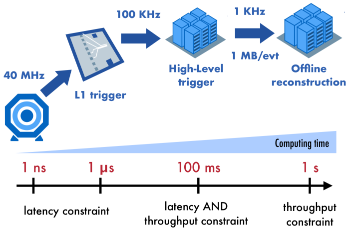

Collecting data —○ DAQ

[Ngadiuba's talk @FastML2020]

Input
48 "trigger cells"
7b floating point
(336b total)

ASIC

Output:
"Super trigger cell" algo
3[16 TC sum] x 13b = 39b

CMS HGCAL cells

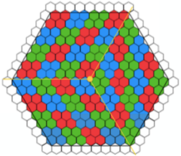

Can we do a better job of encoding the info in those bits w/o so much loss in granularity?

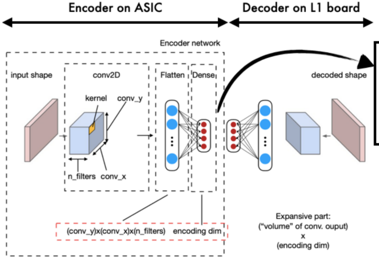

Really need quantized training here to optimize information encoding

Use QKeras!

[More in G. Di Guglielmo talk Monday @ 14:50]

Collecting data

DAQ

Fritzsche's talk @FastML2020

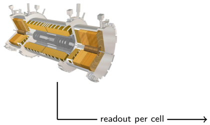

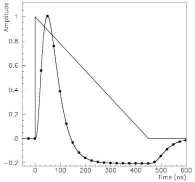

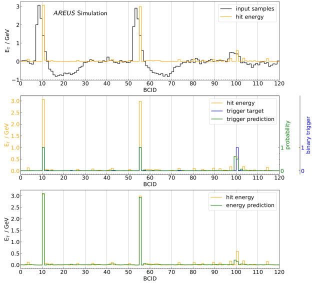

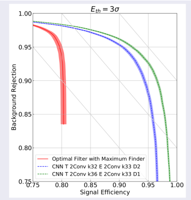

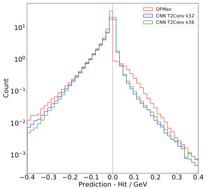

CMS-TDR-021

- Once data is digitalised and pre-selected, it needs to be skimmed more prior to be stored → **trigger**
- Conventionally, a set of selection requirements (cuts)
- **Not efficient** to select interesting events → **use ML models**
- Need to work *really* **fast**, be **lightweight** and programmed to fit into hardware (e.g. **FPGA**)

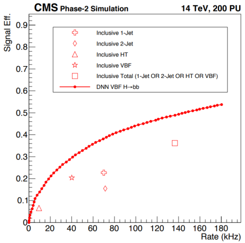

FPGA diagram

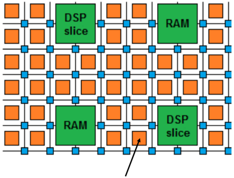

Collecting data —o Trigger

CMS-TDR-021

- Once data is digitalised and pre-selected, it needs to be skimmed more prior to be stored → **trigger**
- Conventionally, a set of selection requirements (cuts)
- **Not efficient** to select interesting events → **use ML models**
- Need to work *really* **fast**, be **lightweight** and programmed to fit into hardware (e.g. **FPGA**)

→ **hls4ml project**
  - deployment of ML models on FPGA with ease
  - quantization and weight pruning
  - improved object reconstruction (muons, taus, tracks)
  - anomaly detection with signal-agnostic triggering

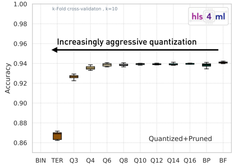

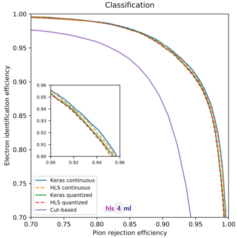

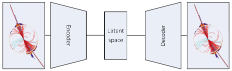

Collecting data → DQM

- After preselection and triggering, we are only interested in "high-quality" data → **data quality monitoring** (DQM)
- can use autoencoder to search for "poor-quality" anomalies

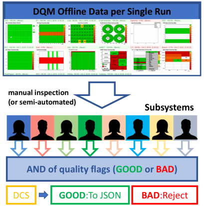

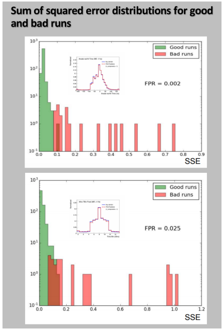

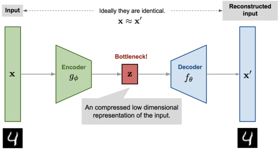

Reconstructing objects

Reconstructing objects → Tracks

Tracking GNN Walk Through

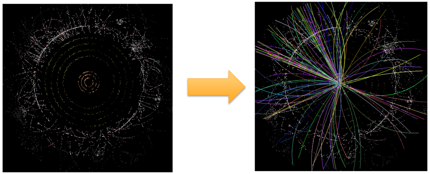

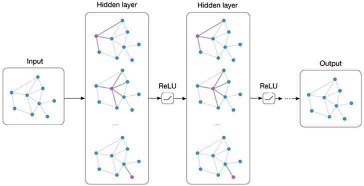

- Traditional Kalman filter will become inefficient with increased luminosity
- It all (seriously) started from TrackML hackathon on Kaggle
- Several research projects emerged, e.g. ACTS, HEP.TrkX and then Exa.TrkX as its follow-up
- Various methods probed (CNN, RNN), finally converged to **graph networks**
- reached $\sim 95\%$ purity/precision

GNN in Particle Physics paper

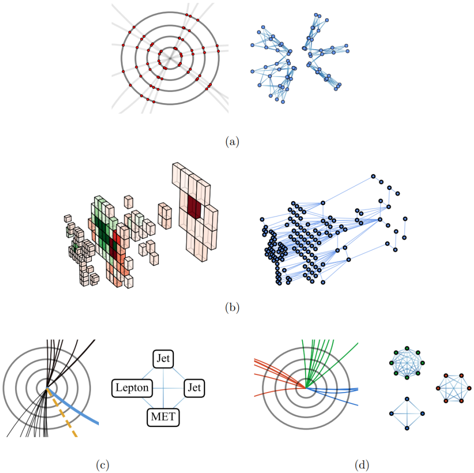

- natural representation
- native treatment of **sparsity**

Qu's talk (CMS restricted)

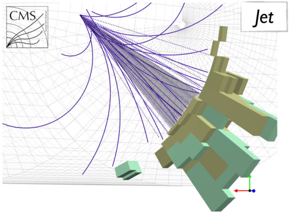

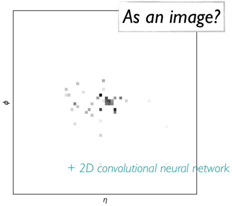

+ 2D convolutional neural network (CNN)

As a Lund plane?

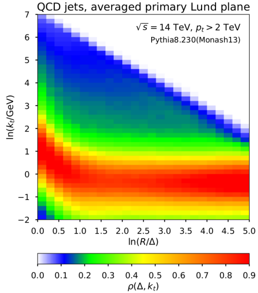

As a sequence?

+ recurrent neural network (RNN)
+ 1D CNN
+ etc.

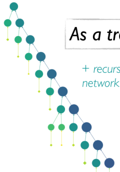

+ recursive neural network (RecNN)

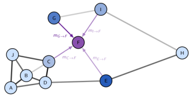

+ graph neural network (GNN)

Reconstructing objects — Jets: Segmentation

You might've heard already about **anti-kt algorithm**, but we can cluster objects differently:

- CV-inspired YOLO + graphs under the hood for cluster reconstruction @LHCb
- GravNet: clustering energy depositions in CMS HGCal with distance-weighted graph nets
- Object condensation: on top of GravNet to to build *one-shot ParticleFlow*

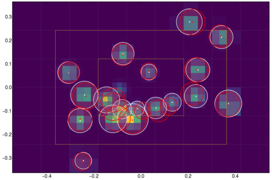

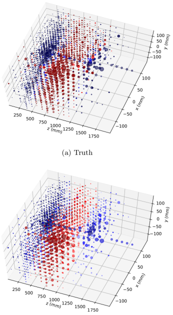

(a) Truth

(b) Reconstructed

Reconstructing objects — Bonus: Neutrino

[Kazuhiro's talk @IML2020]

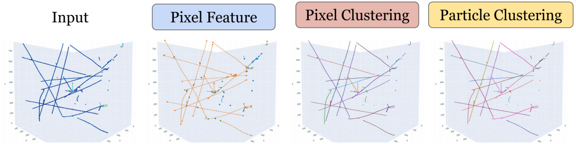

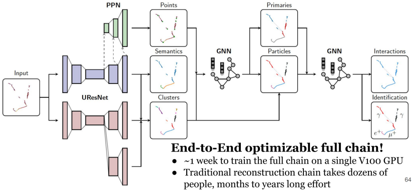

- **b-tagging** [1] , [2]
  - crucial to identify jets from b quark (e.g. $H \to b\bar{b}$)
  - ML brings significant gain ($\sim 50\%$) comparing to classical approaches
  - solid place in production (e.g. DeepJet @CMS)
  - in fact, also used for c/uds/g tagging

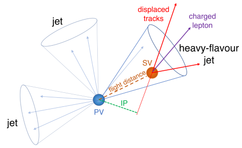

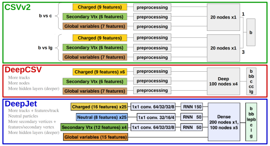

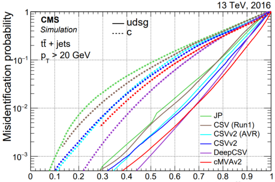

Reconstructing objects — Jets: Tagging

- **b-tagging** [1] , [2]
  - crucial to identify jets from b quark (e.g. $H \to b\bar{b}$)
  - ML brings significant gain ($\sim 50\%$) comparing to classical approaches
  - solid place in production (e.g. DeepJet @CMS)
  - in fact, also used for c/uds/g tagging
- **t-tagging**
  - significant piece in SM puzzle (e.g. ttH, $t\bar{t}$)
  - complex to become benchmarking playground in tagging

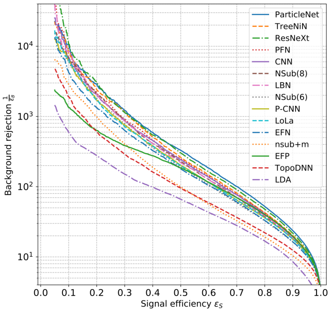

Reconstructing objects — Jets: Cleansing

- **Pile-up mitigation**
  - remove irrelevant contributions from pile-up
  - PUMML: convolutional analysis of LV charged/PU charged/neutral components
  - PuppiML: gated graph-based architecture
  - ABCNet: graph net with attention

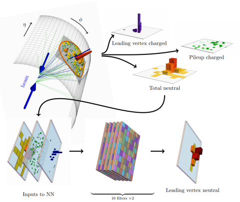

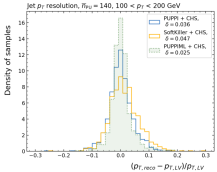

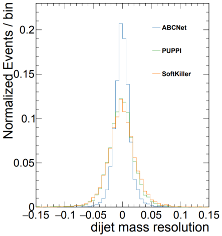

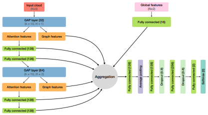

Reconstructing objects — Jets: Cleansing

- **Pile-up mitigation**
  - remove irrelevant contributions from pile-up
  - PUMML: convolutional analysis of LV charged/PU charged/neutral components
  - PuppiML: gated graph-based architecture
  - ABCNet: graph net with attention
- **Super-resolution**
  - enhance the calorimiter's granularity with CNN/GNN
  - able to nicely resolve photons in $\pi_0 \to \gamma\gamma$

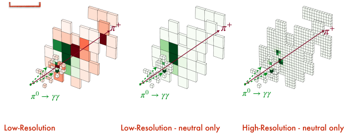

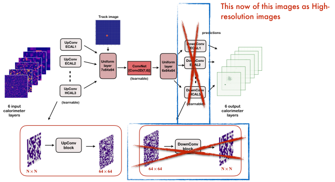

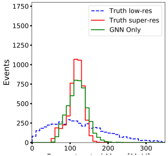

- **Vertices**
  - find ML-reconstructed (secondary) vertices in jets
  - to be further passed to tagging algorithms
- **Leptons**
  - τ: e.g. ABCNet for low-p_T, or general-purpose DeepTau
  - $\mu/e/\gamma$ (were among the first MVA ID)
- **Missing transverse energy** (MET):
  - used to approximate undetected particles' momentum (ν, SUSY)
  - e.g. DeepMET
- **Long-lived particles** (LLP)
  - if you want to search for some SUSY
- **Bosons**
  - tag W/Z/H bosons directly
  - especially in boosted topologies
  - and with a Lund plane

Simulating physics

Simulating physics — Motivation

- What we want to know from data is parameters of the model we study (SM/BSM)
- Traditionally, model's behaviour at the level of observables is obtained by (aka **Monte Carlo**) simulation in the form of **templates** to be fitted to observed data
- **Need to simulate A LOT** of data (e.g. to have low statistical uncertainty of model templates)
- Takes A LOT of time and computational resources
- Can we speed it up?

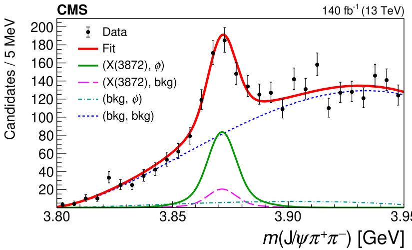

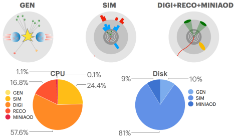

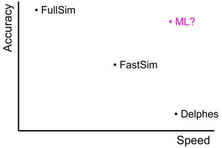

Simulating physics → Motivation

- What we want to know from data is parameters of the model we study (SM/BSM)
- Traditionally, model's behaviour at the level of observables is obtained by (aka **Monte Carlo**) simulation in the form of **templates** to be fitted to observed data
- **Need to simulate A LOT** of data (e.g. to have low statistical uncertainty of model templates)
- Takes A LOT of time and computational resources
- Can we speed it up? → **generative models**

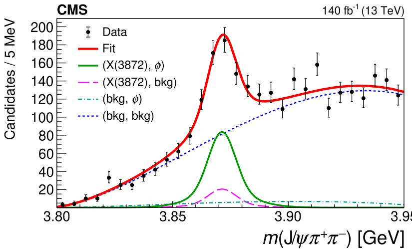

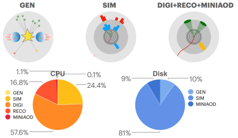

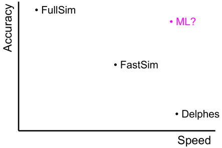

Simulating physics → Examples

- tried images, features, point clouds used as inputs/outputs
- can be *conditional* → specify properties to be generated
- **GAN**
  - historically first: LAGAN , CaloGAN
  - graph-based , end-to-end mapping GEN→RECO levels and bypassing details
    RICH response simulation
- **Autoencoder**
  - e.g. getting high with BIB-AE, or VAE+GAN @LHCb
- crucial to understand and evaluate **systematic uncertainties**
- reached speed-up of $x \lesssim 1000$ comparing to GEANT
- extensive overview here

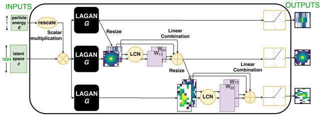

FIG. 4: Composite Generator, illustrating three stream with attentional layer-to-layer dependence.

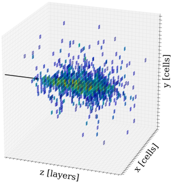

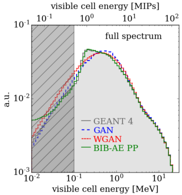

Performing analysis

Performing analysis

Sig vs. Bkg

- need to classify (several) signal processes from (several) backround ones
- afterwards, either apply the cut on model's output and fit observable
- or incorporate it into fitting procedure (e.g. 2D fit)
- **good old classics**, used literally everywhere
- e.g. W^+W^- production (they use Random Forest!!), or $H \to \tau\tau$ CP study @CMS

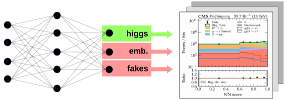

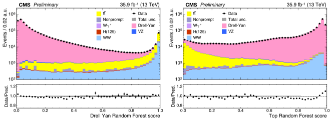

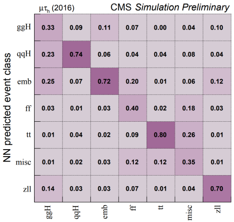

Performing analysis — Anomaly detection

- Normally, we search for new physics having some model in mind
- But if one want to do it **model-independently** – try instead to search for *deviations* $\rightarrow$ **unsupervised learning**
- Also hot topic: ANODE , SALAD

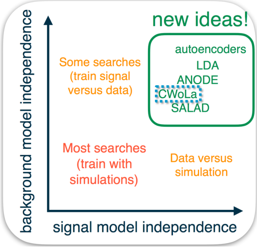

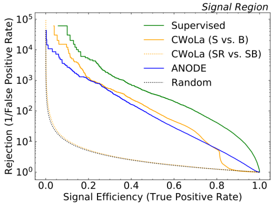

Performing analysis — Anomaly detection

- Normally, we search for new physics having some model in mind
- But if one want to do it **model-independently** – try instead to search for *deviations* $\rightarrow$ **unsupervised learning**
- Also hot topic: ANODE , SALAD
- CWoLa (classification without labels) – used in a (serious) analysis in ATLAS!

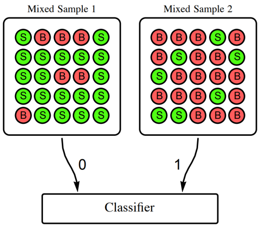

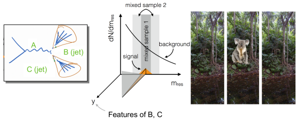

- It was observed, that cutting on classifier output can heavily distort the distribution of observable (e.g. invariant mass) → **mass sculpting**
- This is quite dangerous (not noticed, noticed but reduce sensitivity)
- How to reduce dependency of classifier on some particular variable?
- Mode , DisCo , β-TCVAE-Sensitive , uBoost , adversarial methods

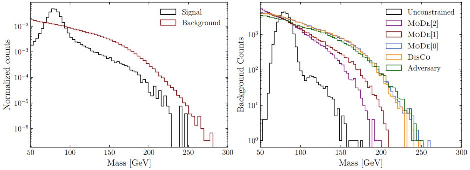

Figure 6. Left: Distributions of signal and background events without selection. **Right:** Background distributions at 50% signal efficiency (true positive rate) for different classifiers. The unconstrained classifier sculpts a peak at the W-boson mass, while other classifiers do not.

Performing analysis — Decorrelation & Adaptation

- It was observed, that cutting on classifier output can heavily distort the distribution of observable (e.g. invariant mass) → **mass sculpting**
- This is quite dangerous (not noticed, noticed but reduce sensitivity)
- How to reduce dependency of classifier on some particular variable?
- Mode , DisCo , β-TCVAE-Sensitive , uBoost , adversarial methods
- Furthermore, we can expand this idea to the **meta-level robustness**: MC/background choice, etc. → domain adaptation

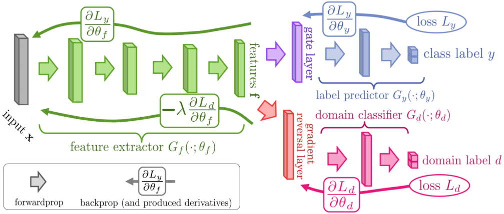

- Every analysis has statistical/systematic uncertainties, so-called **nuisance parameters**
- Vanilla training of ML models doesn't take them into account
- "Optimised" model in reality can be suboptimal – nuisance variations cause large variations of model prediction, thus **diluting sensitivity**
- Moreover, training with e.g. conventional cross-entropy is also suboptimal – we need to **optimise physical metrics**
- Solutions proposed, e.g.:
  - direct optimisation of Asimov Median Significance (AMS)
  - INFERNO, binned Poisson likelihoods
  - likelihood-free inference
- Complex topic and **one of the challenges of ML in HEP** to date!
- see more in this nice review

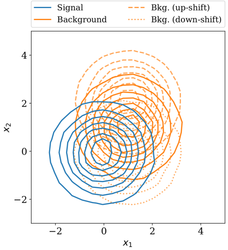

Performing analysis → Nuisances & Inference

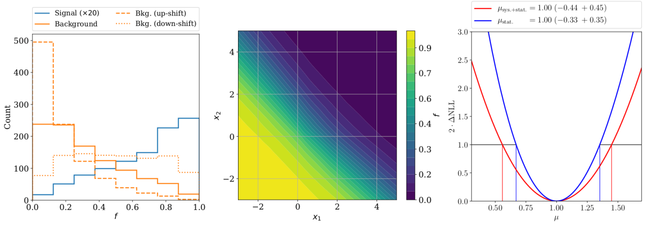

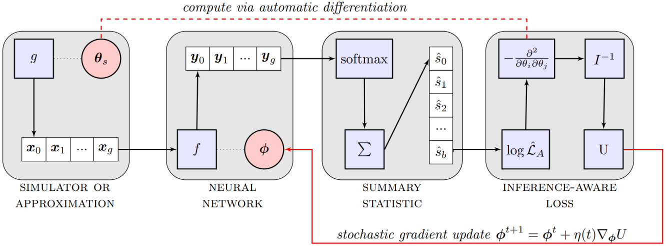

Performing analysis — Interpretability

[Interpretable ML book]

- why ML works? → hand-wavy answer here
  (aka we don't really know)
- but we can try to get insights:
  - bayesian networks to [access per-jet uncertainties]
  - and more bayesian for [accelerators]
  - SHAP values for [feature importance]

DNN

Ensemble

BNN

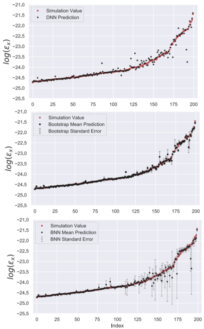

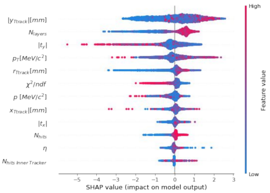

- **Quantum ML**
  - e.g. quantum GANs
  - or quantum classifiers
- **Production**
  - models on your laptop are generally useless for your collegues
  - needs to be implemented into some framework
  - a whole new topic to be aware of

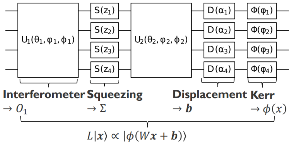

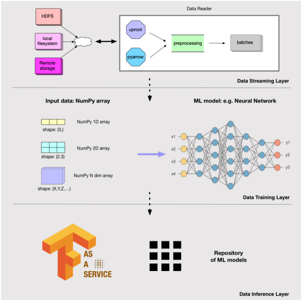

Contemplating future

thoughts from spring 2019

- **research**
  - incorporate physics into the non-physical ML
  - make it (physically) interpretable
  - get closer to ML community (but also be inventive)

thoughts from spring 2019

- **research**
  - incorporate physics into the non-physical ML
  - make it (physically) interpretable
  - get closer to ML community (but also be inventive)
- **production**
  - deploy more models into frameworks in a centralised manner
  - aggregation of hardware (more GPUs to the crowd!)
  - ”make all that cool research available for everyone”

Contemplating future

thoughts from spring 2019

- **research**
  - incorporate physics into the non-physical ML
  - make it (physically) interpretable
  - get closer to ML community (but also be inventive)
- **production**
  - deploy more models into frameworks in a centralised manner
  - aggregation of hardware (more GPUs to the crowd!)
  - ”make all that cool research available for everyone”
- **knowledge**
  - structure all that (exponentionally growing) papers
  - educate people on how to do ML in their analyses
  - ”bookkeeping” of research and production

- **now**
  - "I heard ML is cool, maybe I should use it?"

thoughts from spring 2019

- **now**
  - "I heard ML is cool, maybe I should use it?"
- **soon**
  - *opeping a catalogue of models*
  - "today I feel like going for some boosted-graphy-convo-net sprinkled with sparsed variational tagger"

thoughts from spring 2019

- **now**
  - "I heard ML is cool, maybe I should use it?"
- **soon**
  - *opeping a catalogue of models*
  - "today I feel like going for some boosted-graphy-convo-net sprinkled with sparsed variational tagger"
- **future**
  - "sure, but where's new physics?"

amazing talk on ML&HEP future

Contemplating future

Third wave data scientist

inspired by

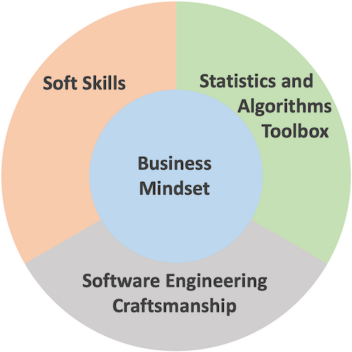

Contemplating future

Third wave data scientist in HEP

inspired by

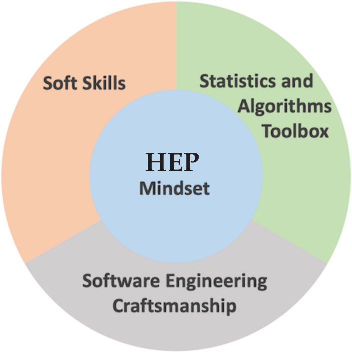

Summary

1. Making theories
2. Building detector
3. Collecting data
4. Reco&ID objects
5. Simulating physics
6. Performing analysis
7. Contemplating future